

Problem Statement: The task is to build a predictive model that can accurately classify mobile phones into predefined price ranges based on various attributes such as battery power, camera features, memory, connectivity options, and more. The dataset provided contains information about several mobile phones, including their specifications and corresponding price ranges.

```
In [1]: #importing libraries
import os
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

sns.set()
%matplotlib inline

import warnings
warnings.filterwarnings("ignore")
```

```
In [2]: dataset = pd.read_csv("train.csv")
dataset.head()
```

```
Out[2]:
```

	battery_power	blue	clock_speed	dual_sim	fc	four_g	int_memory	m_dep	mobile_wt	n_cores	...	px_height	px_width	ram
0	842	0	2.2	0	1	0	7	0.6	188	2	...	20	756	254
1	1021	1	0.5	1	0	1	53	0.7	136	3	...	905	1988	263
2	563	1	0.5	1	2	1	41	0.9	145	5	...	1263	1716	260
3	615	1	2.5	0	0	0	10	0.8	131	6	...	1216	1786	276
4	1821	1	1.2	0	13	1	44	0.6	141	2	...	1208	1212	141

5 rows × 21 columns

```
In [3]: # To avoid data loss
dataset2 = dataset.copy()
```

```
In [4]: #checking information about data
dataset.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2000 entries, 0 to 1999
Data columns (total 21 columns):
#   Column          Non-Null Count  Dtype
---  -
0   battery_power    2000 non-null   int64
1   blue             2000 non-null   int64
2   clock_speed      2000 non-null   float64
3   dual_sim         2000 non-null   int64
4   fc               2000 non-null   int64
5   four_g           2000 non-null   int64
6   int_memory       2000 non-null   int64
7   m_dep            2000 non-null   float64
8   mobile_wt        2000 non-null   int64
9   n_cores          2000 non-null   int64
10  pc               2000 non-null   int64
11  px_height        2000 non-null   int64
12  px_width         2000 non-null   int64
13  ram              2000 non-null   int64
14  sc_h             2000 non-null   int64
15  sc_w             2000 non-null   int64
16  talk_time        2000 non-null   int64
17  three_g          2000 non-null   int64
18  touch_screen     2000 non-null   int64
19  wifi             2000 non-null   int64
20  price_range      2000 non-null   int64
dtypes: float64(2), int64(19)
memory usage: 328.2 KB
```


## EDA

```
In [5]: #statistical analysis
dataset.describe()
```


	battery_power	blue	clock_speed	dual_sim	fc	four_g	int_memory	m_dep	mobile_wt	
<b>count</b>	2000.000000	2000.0000	2000.000000	2000.000000	2000.000000	2000.000000	2000.000000	2000.000000	2000.000000	200
<b>mean</b>	1238.518500	0.4950	1.522250	0.509500	4.309500	0.521500	32.046500	0.501750	140.249000	
<b>std</b>	439.418206	0.5001	0.816004	0.500035	4.341444	0.499662	18.145715	0.288416	35.399655	
<b>min</b>	501.000000	0.0000	0.500000	0.000000	0.000000	0.000000	2.000000	0.100000	80.000000	
<b>25%</b>	851.750000	0.0000	0.700000	0.000000	1.000000	0.000000	16.000000	0.200000	109.000000	
<b>50%</b>	1226.000000	0.0000	1.500000	1.000000	3.000000	1.000000	32.000000	0.500000	141.000000	
<b>75%</b>	1615.250000	1.0000	2.200000	1.000000	7.000000	1.000000	48.000000	0.800000	170.000000	
<b>max</b>	1998.000000	1.0000	3.000000	1.000000	19.000000	1.000000	64.000000	1.000000	200.000000	

8 rows × 21 columns


```
In [6]: sns.jointplot(x='ram',y='price_range',data=dataset,color='red',kind='kde');
```



```
In [7]: #3G support
labels = ["3G-supported", 'Not supported']
values=dataset['three_g'].value_counts().values
fig1, ax1 = plt.subplots()
ax1.pie(values, labels=labels, autopct='%1.1f%%',shadow=True,startangle=90)
plt.show()
```




```
In [8]: #4G support
labels4g = ["4G-supported", 'Not supported']
values4g = dataset['four_g'].value_counts().values
fig1, ax1 = plt.subplots()
ax1.pie(values4g, labels=labels4g, autopct='%1.1f%%', shadow=True, startangle=90)
plt.show()
```




```
In [9]: #battery power vs price_range
sns.boxplot(x="price_range", y="battery_power", data=dataset)
```

```
Out[9]: <AxesSubplot:xlabel='price_range', ylabel='battery_power'>
```



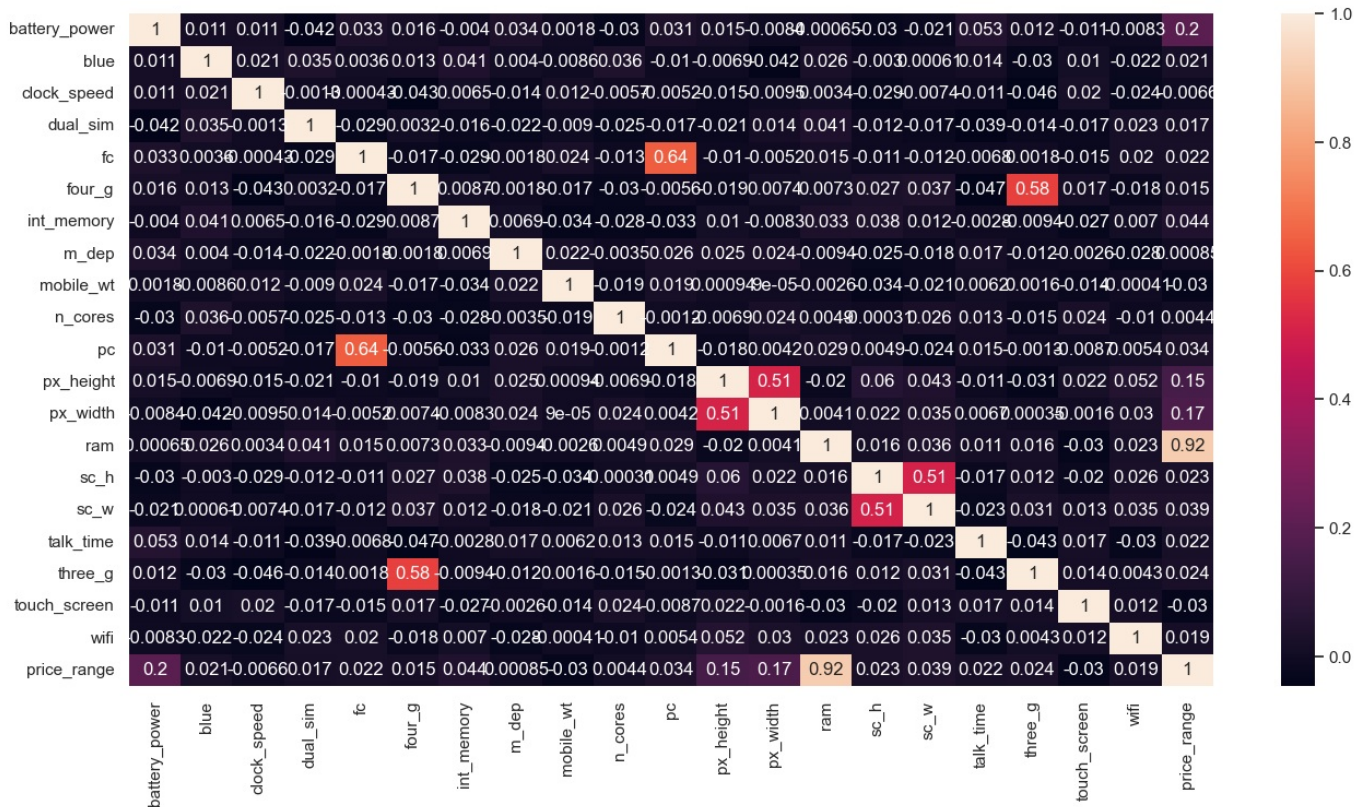
```
In [10]: sns.pointplot(y="talk_time", x="price_range", data=dataset)
```

```
Out[10]: <AxesSubplot:xlabel='price_range', ylabel='talk_time'>
```



```
In [11]: #correlation among columns  
plt.figure(figsize=(16,8))
```

```
sns.heatmap(dataset.corr(), annot = True)
plt.show()
```



## dependent and independent variable

```
In [12]: x=dataset.drop('price_range',axis=1)
```

```
In [13]: x.head()
```

```
Out[13]:
```

	battery_power	blue	clock_speed	dual_sim	fc	four_g	int_memory	m_dep	mobile_wt	n_cores	pc	px_height	px_width	ram	sc_h	sc_w	talk_time	three_g	touch_screen	wifi	price_range
0	842	0	2.2	0	1	0	7	0.6	188	2	2	20	756	25							
1	1021	1	0.5	1	0	1	53	0.7	136	3	6	905	1988	26							
2	563	1	0.5	1	2	1	41	0.9	145	5	6	1263	1716	26							
3	615	1	2.5	0	0	0	10	0.8	131	6	9	1216	1786	27							
4	1821	1	1.2	0	13	1	44	0.6	141	2	14	1208	1212	14							

```
In [14]: y=dataset['price_range']
```

```
In [15]: y.head()
```

```
Out[15]:
```

0	1
1	2
2	2
3	2
4	1

Name: price\_range, dtype: int64

## Train Test split

```
In [16]: from sklearn.model_selection import train_test_split
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.33, random_state=101)
```

## Model

```
In [17]: #logistic regression
from sklearn.linear_model import LogisticRegression
logit_model = LogisticRegression()
logit_model.fit(x_train, y_train)
```

```
Out[17]: LogisticRegression()
```

```
In [18]: #creating function to check models performance
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report
def model_report(model):
    y_pred_train = model.predict(x_train)
    y_pred_test = model.predict(x_test)
    print("Accuracy Score Train:", accuracy_score(y_train,y_pred_train))
    print("Accuracy score test:", accuracy_score(y_test,y_pred_test))
    print("*****" * 9)

    print("Confusion Matrix train:\n ", confusion_matrix(y_train,y_pred_train))
    print("Confusion matrix test: \n", confusion_matrix(y_test,y_pred_test))
    print("*****" * 9)

    print("classification report train: \n", classification_report(y_train,y_pred_train ))
    print("classification report test: \n", classification_report(y_test,y_pred_test ))
```

```
In [19]: model_report(logit_model)
```

```
Accuracy Score Train: 0.6455223880597015
Accuracy score test: 0.6181818181818182
*****
Confusion Matrix train:
[[270  68   4   0]
 [ 57 194  71  26]
 [   1  81 123  96]
 [   0   4  67 278]]
Confusion matrix test:
[[122  35   1   0]
 [ 25  89  32   6]
 [   0  47  78  74]
 [   0   1  31 119]]
*****
classification report train:
              precision    recall  f1-score   support

     0           0.82       0.79       0.81         342
     1           0.56       0.56       0.56         348
     2           0.46       0.41       0.43         301
     3           0.69       0.80       0.74         349

 accuracy                   0.65         1340
 macro avg              0.64       0.64       0.64         1340
 weighted avg           0.64       0.65       0.64         1340

classification report test:
              precision    recall  f1-score   support

     0           0.83       0.77       0.80         158
     1           0.52       0.59       0.55         152
     2           0.55       0.39       0.46         199
     3           0.60       0.79       0.68         151

 accuracy                   0.62         660
 macro avg              0.62       0.63       0.62         660
 weighted avg           0.62       0.62       0.61         660
```

```
In [20]: ## Random Forest Classifier
from sklearn.ensemble import RandomForestClassifier
rfc = RandomForestClassifier()
rfc.fit(x_train, y_train)
```

```
Out[20]: RandomForestClassifier()
```

```
In [21]: model_report(rfc)
```

Accuracy Score Train: 1.0  
 Accuracy score test: 0.8636363636363636  
 \*\*\*\*\*

Confusion Matrix train:

```
[[342  0  0  0]
 [ 0 348  0  0]
 [ 0  0 301  0]
 [ 0  0  0 349]]
```

Confusion matrix test:

```
[[147 11  0  0]
 [13 133  6  0]
 [ 0 26 151 22]
 [ 0  0 12 139]]
```

\*\*\*\*\*

classification report train:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	342
1	1.00	1.00	1.00	348
2	1.00	1.00	1.00	301
3	1.00	1.00	1.00	349
accuracy			1.00	1340
macro avg	1.00	1.00	1.00	1340
weighted avg	1.00	1.00	1.00	1340

classification report test:

	precision	recall	f1-score	support
0	0.92	0.93	0.92	158
1	0.78	0.88	0.83	152
2	0.89	0.76	0.82	199
3	0.86	0.92	0.89	151
accuracy			0.86	660
macro avg	0.86	0.87	0.87	660
weighted avg	0.87	0.86	0.86	660

In [22]: *#above model is overfitted so we need to understand why it happens and then adjust hyperparameters*  
*#hyperparameter tunning*

```
from sklearn.model_selection import GridSearchCV
rf_param_grid = {
    'n_estimators': [100, 200, 15],
    'max_depth': [6, 8, 9],
    'min_samples_split': [2, 5, 10]
}
random_forest_instance = RandomForestClassifier()
rf_grid = GridSearchCV(estimator=random_forest_instance, param_grid=rf_param_grid, cv=5)
rf_grid.fit(x_train, y_train)
print("Best parameters for Random Forest Regressor:", rf_grid.best_params_)
print("Best score for Random Forest Regressor:", rf_grid.best_score_)
```

Best parameters for Random Forest Regressor: {'max\_depth': 9, 'min\_samples\_split': 2, 'n\_estimators': 200}  
 Best score for Random Forest Regressor: 0.864179104477612

In [23]: *#from above code we get a new model of rfc which is rf\_grid, so we will check accuracy for that model*  
 model\_report(rf\_grid)

```
Accuracy Score Train: 0.9992537313432835
Accuracy score test: 0.8621212121212121
*****
```

```
Confusion Matrix train:
```

```
[[342  0  0  0]
 [ 0 347  1  0]
 [ 0  0 301  0]
 [ 0  0  0 349]]
```

```
Confusion matrix test:
```

```
[[150  8  0  0]
 [12 134  6  0]
 [ 0  30 146 23]
 [ 0  0 12 139]]
```

```
*****
```

```
classification report train:
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	342
1	1.00	1.00	1.00	348
2	1.00	1.00	1.00	301
3	1.00	1.00	1.00	349
accuracy			1.00	1340
macro avg	1.00	1.00	1.00	1340
weighted avg	1.00	1.00	1.00	1340

```
classification report test:
```

	precision	recall	f1-score	support
0	0.93	0.95	0.94	158
1	0.78	0.88	0.83	152
2	0.89	0.73	0.80	199
3	0.86	0.92	0.89	151
accuracy			0.86	660
macro avg	0.86	0.87	0.86	660
weighted avg	0.87	0.86	0.86	660

```
In [24]: from sklearn.neighbors import KNeighborsClassifier
knn = KNeighborsClassifier(n_neighbors=10)
knn.fit(x_train,y_train)
```

```
Out[24]: KNeighborsClassifier(n_neighbors=10)
```

```
In [25]: model_report(knn)
```



Accuracy Score Train: 0.9462686567164179  
 Accuracy score test: 0.9212121212121213  
 \*\*\*\*\*

Confusion Matrix train:

```
[[334  8  0  0]
 [ 14 331  3  0]
 [  0 18 271 12]
 [  0  0 17 332]]
```

Confusion matrix test:

```
[[157  1  0  0]
 [  6 143  3  0]
 [  0 20 169 10]
 [  0  0 12 139]]
```

\*\*\*\*\*

classification report train:

	precision	recall	f1-score	support
0	0.96	0.98	0.97	342
1	0.93	0.95	0.94	348
2	0.93	0.90	0.92	301
3	0.97	0.95	0.96	349
accuracy			0.95	1340
macro avg	0.95	0.94	0.95	1340
weighted avg	0.95	0.95	0.95	1340

classification report test:

	precision	recall	f1-score	support
0	0.96	0.99	0.98	158
1	0.87	0.94	0.91	152
2	0.92	0.85	0.88	199
3	0.93	0.92	0.93	151
accuracy			0.92	660
macro avg	0.92	0.93	0.92	660
weighted avg	0.92	0.92	0.92	660

In [26]: *#using elbow method to get best values of hyperparameter- n\_neighbors*

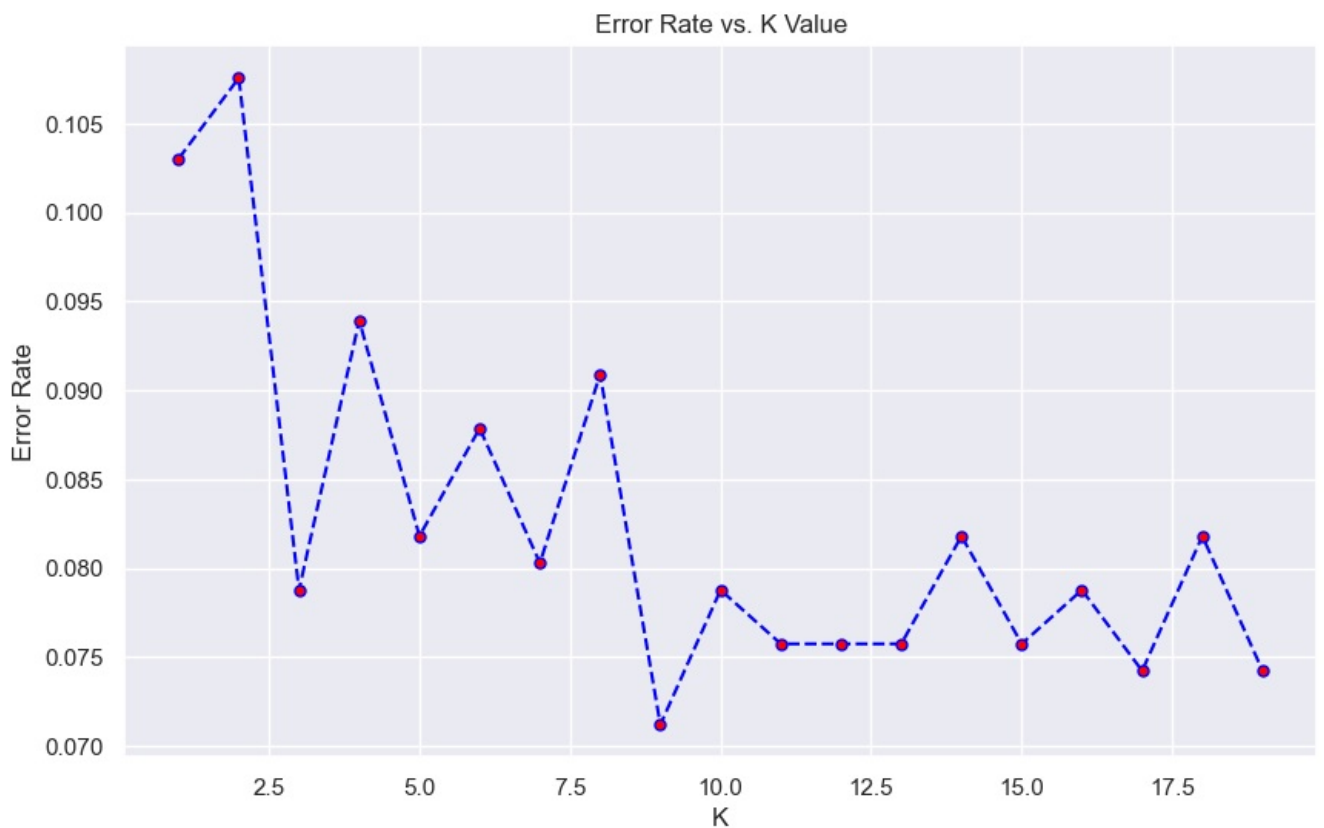
```
error_rate = []
for i in range(1,20):

    knn = KNeighborsClassifier(n_neighbors=i)
    knn.fit(x_train,y_train)
    pred_i = knn.predict(x_test)
    error_rate.append(np.mean(pred_i != y_test))
```

In [27]: *#lets visualize these neighbors to get clear picture*

```
plt.figure(figsize=(10,6))
plt.plot(range(1,20),error_rate,color='blue', linestyle='dashed', marker='o',
        markerfacecolor='red', markersize=5)
plt.title('Error Rate vs. K Value')
plt.xlabel('K')
plt.ylabel('Error Rate')
```

Out[27]: Text(0, 0.5, 'Error Rate')



Result is KNN model performs best by comparing their accuracies

```
In [28]: test=pd.read_csv('test.csv')
test.head()
```

```
Out[28]:
```

	id	battery_power	blue	clock_speed	dual_sim	fc	four_g	int_memory	m_dep	mobile_wt	...	pc	px_height	px_width	ram
0	1	1043	1	1.8	1	14	0	5	0.1	193	...	16	226	1412	3476
1	2	841	1	0.5	1	4	1	61	0.8	191	...	12	746	857	3895
2	3	1807	1	2.8	0	1	0	27	0.9	186	...	4	1270	1366	2396
3	4	1546	0	0.5	1	18	1	25	0.5	96	...	20	295	1752	3893
4	5	1434	0	1.4	0	11	1	49	0.5	108	...	18	749	810	1773

5 rows × 21 columns

```
In [29]: test=test.drop('id',axis=1)
test.head()
```

Out[29]:	battery_power	blue	clock_speed	dual_sim	fc	four_g	int_memory	m_dep	mobile_wt	n_cores	pc	px_height	px_width	ra
0	1043	1	1.8	1	14	0	5	0.1	193	3	16	226	1412	34
1	841	1	0.5	1	4	1	61	0.8	191	5	12	746	857	38
2	1807	1	2.8	0	1	0	27	0.9	186	3	4	1270	1366	23
3	1546	0	0.5	1	18	1	25	0.5	96	8	20	295	1752	38
4	1434	0	1.4	0	11	1	49	0.5	108	6	18	749	810	17

```
In [30]: # lets predict price by feeding data to best model i.e., KNN
predict_price = knn.predict(test)
```

```
In [31]: predict_price
```

```
Out[31]: array([3, 3, 2, 3, 1, 3, 3, 1, 3, 0, 3, 3, 0, 0, 2, 0, 2, 1, 3, 2, 1, 3,
1, 1, 3, 0, 2, 0, 3, 0, 2, 0, 3, 0, 0, 1, 3, 1, 2, 1, 1, 2, 0, 0,
0, 1, 0, 3, 1, 2, 1, 0, 3, 0, 3, 1, 3, 1, 1, 3, 3, 2, 0, 1, 1, 1,
2, 3, 1, 2, 1, 2, 2, 3, 3, 0, 2, 0, 2, 3, 0, 3, 3, 0, 3, 0, 3, 1,
3, 0, 1, 2, 2, 1, 2, 1, 0, 2, 1, 2, 1, 0, 0, 3, 0, 2, 0, 1, 2, 3,
3, 3, 1, 3, 3, 3, 3, 1, 3, 0, 0, 3, 2, 1, 2, 0, 3, 2, 3, 1, 0, 2,
1, 1, 3, 1, 1, 0, 3, 2, 1, 3, 1, 3, 2, 3, 3, 3, 2, 3, 2, 3, 1, 0,
3, 2, 3, 3, 3, 3, 2, 2, 3, 3, 3, 3, 1, 0, 3, 0, 0, 0, 2, 1, 0, 1,
0, 0, 1, 2, 1, 0, 0, 1, 1, 2, 2, 1, 0, 0, 0, 0, 0, 3, 1, 0, 2, 2,
3, 3, 1, 1, 3, 3, 3, 2, 2, 1, 1, 0, 1, 2, 0, 2, 3, 3, 0, 2, 0, 3,
2, 3, 3, 1, 0, 1, 0, 3, 0, 1, 0, 2, 2, 1, 2, 0, 3, 0, 3, 1, 2, 0,
0, 2, 1, 3, 3, 3, 1, 1, 3, 0, 0, 2, 3, 3, 1, 3, 1, 1, 3, 2, 1, 2,
3, 3, 3, 1, 0, 1, 2, 3, 1, 1, 3, 2, 0, 3, 0, 0, 2, 0, 0, 3, 2, 3,
3, 2, 1, 3, 3, 2, 3, 1, 2, 1, 2, 0, 2, 3, 1, 0, 0, 3, 0, 3, 0, 1,
2, 0, 2, 3, 1, 3, 2, 2, 1, 2, 0, 0, 0, 1, 3, 2, 0, 0, 0, 3, 2, 0,
2, 3, 1, 2, 2, 3, 1, 3, 3, 1, 2, 3, 3, 3, 0, 3, 0, 3, 1, 3, 1,
2, 3, 0, 1, 0, 3, 1, 3, 1, 3, 0, 0, 0, 0, 2, 0, 0, 2, 1, 1, 2, 2,
2, 0, 1, 0, 0, 3, 2, 0, 3, 1, 2, 2, 1, 2, 3, 1, 1, 2, 2, 1, 2, 0,
1, 1, 0, 3, 2, 1, 0, 1, 0, 0, 1, 1, 0, 0, 0, 2, 2, 3, 2, 3, 0, 3,
0, 3, 0, 1, 1, 1, 2, 0, 3, 2, 3, 3, 1, 3, 1, 3, 1, 3, 2, 0, 1, 2,
1, 1, 0, 0, 0, 1, 2, 1, 0, 3, 2, 0, 2, 2, 0, 0, 3, 1, 1, 0, 3, 2,
3, 0, 3, 0, 2, 3, 3, 3, 0, 2, 0, 2, 3, 0, 1, 1, 0, 0, 1, 1, 1, 3,
3, 3, 2, 3, 1, 1, 2, 3, 3, 3, 2, 0, 2, 1, 2, 2, 1, 0, 2, 2, 0, 0,
0, 3, 1, 0, 2, 2, 2, 0, 3, 0, 2, 2, 0, 3, 0, 2, 3, 0, 1, 1, 3, 3,
1, 1, 1, 3, 2, 0, 3, 1, 2, 0, 3, 3, 1, 2, 2, 2, 3, 0, 1, 2, 3, 1,
3, 2, 3, 1, 1, 1, 0, 3, 1, 0, 3, 2, 3, 2, 0, 3, 3, 3, 2, 3, 3, 1,
2, 0, 2, 3, 3, 1, 0, 1, 1, 2, 2, 2, 0, 0, 2, 2, 3, 2, 0, 2, 1, 3,
3, 0, 1, 3, 0, 2, 1, 1, 0, 0, 2, 1, 0, 1, 1, 2, 2, 0, 2, 2, 1, 0,
3, 0, 0, 3, 2, 0, 0, 0, 0, 0, 3, 0, 3, 1, 3, 1, 1, 3, 3, 0, 1, 1,
3, 2, 3, 2, 0, 3, 0, 2, 0, 2, 0, 0, 1, 1, 1, 2, 1, 3, 1, 3, 2, 2,
1, 3, 2, 0, 1, 2, 0, 3, 3, 0, 2, 1, 1, 2, 0, 3, 2, 0, 3, 2, 3, 0,
0, 3, 0, 2, 2, 3, 2, 2, 2, 1, 1, 3, 0, 1, 0, 1, 2, 1, 0, 0, 1,
0, 0, 3, 0, 1, 2, 0, 0, 1, 1, 3, 0, 3, 2, 3, 0, 0, 1, 2, 2, 1, 0,
1, 2, 0, 1, 1, 0, 0, 3, 3, 0, 3, 1, 2, 3, 0, 1, 0, 2, 2, 0, 3, 1,
0, 3, 0, 1, 0, 3, 3, 3, 2, 3, 0, 3, 2, 0, 0, 0, 2, 3, 2, 0, 1, 1,
2, 1, 0, 3, 2, 0, 3, 1, 2, 1, 1, 1, 3, 1, 1, 1, 2, 1, 0, 1, 2, 0,
3, 0, 0, 0, 0, 2, 3, 3, 3, 0, 1, 2, 1, 1, 0, 0, 2, 1, 0, 2, 0, 3,
2, 2, 1, 2, 0, 2, 1, 3, 0, 0, 3, 2, 3, 0, 0, 2, 3, 3, 1, 2, 2, 1,
0, 0, 2, 3, 0, 3, 0, 0, 0, 2, 2, 1, 2, 0, 3, 2, 1, 2, 3, 3, 0, 1,
1, 2, 1, 2, 2, 0, 1, 3, 1, 1, 3, 1, 2, 3, 1, 1, 1, 1, 3, 2, 0, 2,
3, 0, 2, 3, 2, 2, 2, 3, 2, 0, 1, 2, 0, 2, 1, 1, 2, 2, 2, 1, 2, 1,
0, 1, 3, 1, 0, 1, 2, 3, 1, 0, 0, 3, 2, 2, 3, 0, 3, 3, 2, 1, 3, 0,
1, 3, 1, 1, 1, 3, 2, 0, 3, 0, 2, 3, 0, 3, 2, 2, 3, 1, 0, 2, 3,
1, 0, 2, 1, 2, 1, 2, 0, 2, 2, 0, 2, 3, 2, 3, 0, 2, 1, 1, 2, 2, 3,
3, 0, 2, 1, 2, 1, 3, 0, 1, 3, 0, 1, 0, 0, 3, 3, 2, 0, 0, 0, 0, 3,
2, 3, 3, 0, 0, 2, 1, 0, 2, 2], dtype=int64)
```

```
In [32]: #let's concatenate with test data
test['price_range']=predict_price
```

```
In [33]: test.head()
```

Out[33]:	battery_power	blue	clock_speed	dual_sim	fc	four_g	int_memory	m_dep	mobile_wt	n_cores	...	px_height	px_width	ra
0	1043	1	1.8	1	14	0	5	0.1	193	3	...	226	1412	347
1	841	1	0.5	1	4	1	61	0.8	191	5	...	746	857	385
2	1807	1	2.8	0	1	0	27	0.9	186	3	...	1270	1366	235
3	1546	0	0.5	1	18	1	25	0.5	96	8	...	295	1752	385
4	1434	0	1.4	0	11	1	49	0.5	108	6	...	749	810	177

5 rows × 21 columns

In [ ]:

In [ ]:

In [ ]:

Loading [MathJax]/jax/output/CommonHTML/fonts/TeX/fontdata.js